

Deep Neural Networks for Natural Language Processing (AI6127)

JUNG-JAE KIM

TUTORIAL 1: TF-IDF

Question 1: Consider these documents

- Doc1 breakthrough drug for schizophrenia
 - Doc2 new schizophrenia drug
 - Doc3 new approach for treatment of schizophrenia
 - Doc4 new hopes for schizophrenia patients
- Draw a term-document count matrix for the document collection
 - Calculate the inverse document frequencies of the terms in the document collection
 - Draw a term-document TF-IDF weighted matrix for the document collection
 - Find the document in the collection that is the closest to the new document “schizophrenia drug” by using cosine similarity between the TF-IDF vectors of the documents

Term-document count matrices

- Consider the number of occurrences of a word in a document:
 - Each document is a count vector in $\mathbb{N}^{|V|}$: a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

Inverse document frequency (idf)

- Rare terms are more informative than frequent terms
 - Stop words (e.g. the, of)
- df_t is the document frequency of t : the number of documents in training data that contain t
 - df_t is an inverse measure of the informativeness of t
 - $df_t \leq N$
- We define the idf (inverse document frequency) of t by
$$idf_t = \log_{10} (N/df_t)$$
 - We use $\log (N/df_t)$ instead of N/df_t to “dampen” the effect of idf.

The base of the log is not important.

tf-idf vector

- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = (1 + \log \text{tf}_{t,d}) \times \log_{10}(N / \text{df}_t)$$

- Increases with the number of occurrences within a document
- Increases with the rarity of the term in the collection

Document similarity: Cosine similarity

- E.g. Find the document (d) in given collection that is the closest to a new document (q)
 - \vec{d} and \vec{q} are TF-IDF vectors

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \cdot \frac{\vec{d}}{|\vec{d}|} = \frac{\sum_{i=1}^{|V|} q_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

Diagram annotations:

- A box labeled "Dot product" points to the $\vec{q} \cdot \vec{d}$ term in the first fraction.
- A box labeled "Unit vectors" points to the $\frac{\vec{q}}{|\vec{q}|}$ and $\frac{\vec{d}}{|\vec{d}|}$ terms in the second fraction.

Question 2: Read and understand the example implementation of sentiment analysis of movie reviews by using binary or TF-IDF vectors

- <https://drive.google.com/file/d/1DyDKHSg6yTpDNQgv5bMfZKi-VDvCHVkJ/view?usp=sharing>

Hands-on

Answer 1: Draw a term-document count matrix for the document collection

	Doc1	Doc2	Doc3	Doc4
Approach	0	0	1	0
Breakthrough	1	0	0	0
Drug	1	1	0	0
For	1	0	1	1
Hopes	0	0	0	1
New	0	1	1	1
Of	0	0	1	0
Patients	0	0	0	1
Schizophrenia	1	1	1	1
treatment	0	0	1	0

Doc1 breakthrough drug for schizophrenia

Doc2 new schizophrenia drug

Doc3 new approach for treatment of schizophrenia

Doc4 new hopes for schizophrenia patients

Answer 1: Calculate the inverse document frequencies of the terms in the document collection

	DF	IDF
Approach	1	0.60
Breakthrough	1	0.60
Drug	2	0.30
For	3	0.12
Hopes	1	0.60
New	3	0.12
Of	1	0.60
Patients	1	0.60
Schizophrenia	4	0.00
treatment	1	0.60

Doc1 breakthrough drug for schizophrenia

Doc2 new schizophrenia drug

Doc3 new approach for treatment of schizophrenia

Doc4 new hopes for schizophrenia patients

$$\text{idf}_t = \log_{10} (N/\text{df}_t)$$

Answer 1: Draw a term-document TF-IDF weighted matrix for the document collection

	Doc1	Doc2	Doc3	Doc4
Approach	0	0	0.60	0
Breakthrough	0.60	0	0	0
Drug	0.30	0.30	0	0
For	0.12	0	0.12	0.12
Hopes	0	0	0	0.60
New	0	0.12	0.12	0.12
Of	0	0	0.60	0
Patients	0	0	0	0.60
Schizophrenia	0.00	0.00	0.00	0.00
treatment	0	0	0.60	0

$$w_{t,d} = (1 + \log \text{tf}_{t,d}) \times \log_{10}(N / \text{df}_t)$$

Answer 1: Find the document in the collection that is the closest to the new document “schizophrenia drug”

	Doc1	Doc2	Doc3	Doc4	Doc5
Approach	0	0	0.60	0	0
Breakthrough	0.60	0	0	0	0
Drug	0.30	0.30	0	0	0.30
For	0.12	0	0.12	0.12	0
Hopes	0	0	0	0.60	0
New	0	0.12	0.12	0.12	0
Of	0	0	0.60	0	0
Patients	0	0	0	0.60	0
Schizophrenia	0.00	0.00	0.00	0.00	0.00
treatment	0	0	0.60	0	0
Length	0.68	0.32	1.05	0.87	0.30

	Doc1	Doc2	Doc3	Doc4	Doc5
Approach	0	0	0.57	0	0
Breakthrough	0.88	0	0	0	0
Drug	0.44	0.93	0	0	1
For	0.18	0	0.11	0.14	0
Hopes	0	0	0	0.69	0
New	0	0.37	0.11	0.14	0
Of	0	0	0.57	0	0
Patients	0	0	0	0.69	0
Schizophrenia	0.00	0.00	0.00	0.00	0.00
treatment	0	0	0.57	0	0

Answer 1: Find the document in the collection that is the closest to the new document “schizophrenia drug”

$$\cos(\text{Doc2}, \text{Doc5}) = 0.93$$

$$\begin{aligned}\cos(\vec{q}, \vec{d}) &= \frac{\vec{q} \cdot \vec{d}}{|\vec{q}| |\vec{d}|} = \frac{\vec{q}}{|\vec{q}|} \cdot \frac{\vec{d}}{|\vec{d}|} \\ &= \frac{\sum_{i=1}^{|V|} q_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2}} \frac{\sum_{i=1}^{|V|} d_i}{\sqrt{\sum_{i=1}^{|V|} d_i^2}}\end{aligned}$$

	Doc1	Doc2	Doc3	Doc4	Doc5
Approach	0	0	0.57	0	0
Breakthrough	0.88	0	0	0	0
Drug	0.44	0.93	0	0	1
For	0.18	0	0.11	0.14	0
Hopes	0	0	0	0.69	0
New	0	0.37	0.11	0.14	0
Of	0	0	0.57	0	0
Patients	0	0	0	0.69	0
Schizophrenia	0.00	0.00	0.00	0.00	0.00
treatment	0	0	0.57	0	0

Answer 2: Download and import movie review dataset

```
import nltk
```

```
nltk.download('movie_reviews')
```

```
from nltk.corpus import movie_reviews
```

Answer 2: Data Processing in Boolean format

```
documents =  
[(list(movie_reviews.words(fileid)),  
category) \n  
    for category in  
movie_reviews.categories() \n  
    for fileid in  
movie_reviews.fileids(category)]  
# print out the size of documents  
print(len(documents))  
  
# shuffle documents, reorganize the order
```

```
of documents  
random.shuffle(documents)  
  
# split documents into training and test  
data  
doc_train, doc_test = documents[100:],  
documents[:100]  
  
# print out the size of training and test  
documents  
print(len(doc_train), len(doc_test))
```

Answer 2: Category set and examples

display category set

```
print(doc_train[0][1])
```

```
categories = set([c for d, c in  
doc_train])
```

```
print(categories)
```

display example movie review and
its category

```
print(doc_train[0][0])
```


Answer 2: Vocabulary

get the word frequency distribution of training data

```
all_words = nltk.FreqDist(w.lower() for d, c in doc_train for w in d)
```

get the top-2000 most common words

```
word_features = [word for word, num in all_words.most_common(2000)]
```

display first 20 words in word_features

```
print(word_features[:20])
```

display the size of all_words

```
all_word_num = len(all_words)
```

```
print(all_word_num)
```

display top-k most common words, k in [0, 50) is generated randomly

```
k = random.randint(0, 50)
```

```
print(k)
```

```
all_words.pprint(k)
```

plot the word frequency distribution of top-k most common words

```
all_words.plot(k, title='word distribution')
```

Answer 2: Binary vectors

define a function, return a dict, the type of key is str and value is Boolean

```
def document_features(document):
```

```
    document_words = set(document)
```

```
    features = {}
```

```
    for word in word_features:
```

```
        features[word] = (word in document_words)
```

```
    return features
```

convert documents into Boolean vectors

```
train_set = [(document_features(d), c) for (d,c) in doc_train]
```

```
test_set = [(document_features(d), c) for (d,c) in doc_test]
```

print example features

```
print(train_set[0][0])
```

Answer 2: Training SVM with binary vectors

```
# train an SVM model with training set
import nltk.classify
from sklearn.svm import LinearSVC

# train the LinearSVC classifier using training data
classifier.train(train_set)

# wrapper sklearn classifier using nltk
classifier =
nltk.classify.SklearnClassifier(LinearSVC(
))

# apply the trained SVM model for test set
accuracy =
nltk.classify.accuracy(classifier, test_set)
print(accuracy)
```

Answer 2: Training Random Forest with binary vectors

train an RF model with training set

```
from sklearn.ensemble import  
RandomForestClassifier
```

```
classifier =  
nltk.classify.SklearnClassifier(Rando  
mForestClassifier())  
  
classifier.train(train_set)
```

apply the trained RF model for test set

```
accuracy =  
nltk.classify.accuracy(classifier,  
test_set)  
  
print(accuracy)
```

Answer 2: TF-IDF vectors

convert documents into tfidf vectors

```
from sklearn.feature_extraction.text  
import TfidfVectorizer
```

```
tfidf =  
TfidfVectorizer(stop_words='english')
```

```
train_text = [' '.join(d) for (d,c) in  
doc_train]
```

```
train_tfidf =  
tfidf.fit_transform(train_text)  
train_cats = [c for d, c in doc_train]
```

```
test_text = [' '.join(d) for (d,c) in  
doc_test]
```

```
test_tfidf = tfidf.transform(test_text)  
test_cats = [c for d, c in doc_test]
```

Answer 2: Training Random Forest with TF-IDF vectors

```
# train an RF model with training set print(accuracy)
```

```
classifier = RandomForestClassifier()
```

```
classifier.fit(train_tfidf, train_cats)
```

```
# apply the trained RF model for test  
set
```

```
accuracy = classifier.score(test_tfidf,  
test_cats)
```