

1. Read and understand the example implementation of n-gram language modeling at https://pytorch.org/tutorials/beginner/nlp/word_embeddings_tutorial.html
 - a. An n-gram is a chunk of n consecutive words.
 - i. unigrams: "the", "students", "opened", "their"
 - ii. bigrams: "the students", "students opened", "opened their"
 - iii. trigrams: "the students opened", "students opened their"
 - b. Language Modeling is the task of predicting what word comes next. Formally, given a sequence of words $x^{(1)}, x^{(2)}, \dots, x^{(t)}$, compute the probability distribution of the next word $x^{(t+1)}$: $p(x^{(t+1)} | x^{(t)}, \dots, x^{(1)})$
 - i. where $x^{(t+1)}$ can be any word in the vocabulary $V = \{w_1, \dots, w_{|V|}\}$
2. Give answer codes for the exercise "Computing Word Embeddings: Continuous Bag-of-Words" based on the example implementation of n-gram language modeling, as follows:
 - a. Filling up the `__init__` and forward functions
 - b. Giving codes for training and for displaying word vectors of selected words
3. Run the notebook "Gensim word vector visualization of various word vectors" at <http://web.stanford.edu/class/cs224n/materials/Gensim.zip>
 - a. Changing the datapath as follows:

```
!wget http://nlp.stanford.edu/data/glove.6B.zip
!unzip glove.6B.zip
glove_file = datapath('/content/glove.6B.100d.txt')
```