

# Deep Neural Networks for Natural Language Processing (AI6127)

---

JUNG-JAE KIM

TUTORIAL 4: CONVOLUTIONAL NEURAL NETWORKS

# Question 1

---

- Run the iPython notebook (Section 5.3, “Document classification with CNN”) of the following site on your machine:
  - <https://github.com/joosthub/PyTorchNLPBook> (Repository of source codes explained in the book “Natural Language Processing with PyTorch” by Delip Rao and Brian McMahan)
  - It requires data file downloading and data pre-processing

# Hint 1

---

- Download ag\_news dataset (see data/get-all-data.sh)
  - `cd data; python download.py 1hjAZJJVyez-tjaUSwQyMBMVbW68Kgyzn ./ag_news/news.csv`
- Download Glove word embeddings
  - <http://nlp.stanford.edu/data/glove.6B.zip>: unzip and place it under data/glove
- Run 5\_3\_Munging\_AG\_News notebook
- Run 5\_3\_Document\_Classification\_with\_CNN notebook
  - Fix `load_glove_from_file` function in section “general utilities”
    - with `open(glove_filepath, "r", encoding='utf8')` as `fp`:

# Question 2: Answer the following questions about the notebook

---

- How many layers of 1D CNN are used in the model?
  - Hint: See Section “The Model: NewsClassifier”, `__init__` function
- What other layers are included in the model?
  - Hint: See Section “The Model: NewsClassifier”, `__init__` and forward functions
- How many channels does each 1D CNN produce as output?
  - Hints:
    - See Section “The Model: NewsClassifier”, `__init__` function
    - Sections “Initializations” and “Settings and some prep work”

# Question 2: Answer the following questions about the notebook

---

- What is difference between the padding introduced in the lecture slides and the padding introduced in the notebook?
  - Hints:
    - Section “The vectorizer”, `vectorize` function
    - Section “The dataset”, `__init__` and `__getitem__` functions
    - Section “The Vocabulary”, `__init__` function

# A 1D convolution for text with padding = 1

∅	0.0	0.0	0.0	0.0
tentative	0.2	0.1	-0.3	0.4
deal	0.5	0.2	-0.3	-0.1
reached	-0.2	-0.3	-0.2	0.4
to	0.3	-0.3	0.1	0.1
keep	0.2	-0.3	0.4	0.2
government	0.1	0.2	-0.1	-0.1
open	-0.4	-0.4	0.2	0.3
∅	0.0	0.0	0.0	0.0

∅,t,d	-0.6
t,d,r	-1.1
d,r,t	-0.4
r,t,k	-3.9
t,k,g	-0.2
k,g,o	0.3
g,o, ∅	-0.5

3	1	2	-3
-1	2	1	-3
1	1	-1	1

Apply a **filter**  
of size 3

# Question 2: Answer the following questions about the notebook

---

- What pooling method is used in the notebook?
  - Hint: See Section “The Model: NewsClassifier”, forward function
- Which regularisation method is used in the notebook?
  - Hint: See Section “The Model: NewsClassifier”, `__init__` and forward functions

# Regularization: dropout

---


- Create masking vector  $r$  of Bernoulli random variables with probability  $p$  (a hyperparameter) of being 1
- Delete features during training:  $y = \text{softmax}(W^{(S)}(r \circ z) + b)$ 
  - E.g. randomly set 50% of the inputs to each neuron to 0, at each instance
- Reasoning: Prevents co-adaptation (overfitting to seeing specific feature constellations)
  - See the link below for details
- At test time, no dropout, scale down final vector by probability  $p$ :  
$$\hat{W}^{(S)} = pW^{(S)}$$



# Regularization: $l_2$ norm

- Recall:  $\|v\|_2 = \sqrt{v_1^2 + v_2^2 + v_3^2 + \dots}$
- Constrain  $l_2$  norms of weight vectors of each class (row in softmax weight  $W^{(s)}$ ) to be upper-bounded by a fixed number  $s$  (also a hyperparameter)
  - If  $\|W_c^{(s)}\|_2 > s$ , then rescale it so that:  $\|W_c^{(s)}\|_2 = s$
- Or, add regularization to loss function
  - Helps to prevent **overfitting** when we have a lot of features

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N -\log \left( \frac{e^{f_{y_i}}}{\sum_{c=1}^C e^{f_c}} \right) + \lambda \sum_k \theta_k^2$$



# Hands-on

---

## Question 2: Answer the following questions about the notebook

---

- Question 2.1: How many layers of 1D CNN are used in the model?
- Hint: See Section “The Model: NewsClassifier”, `__init__` function

# Answer 2.1

---

- Section “The Model: NewsClassifier”

- `__init__` function

```
self.convnet = nn.Sequential(  
    nn.Conv1d(in_channels=embedding_size,  
              out_channels=num_channels, kernel_size=3),  
    nn.ELU(),  
    nn.Conv1d(in_channels=num_channels, out_channels=num_channels,  
              kernel_size=3, stride=2),  
    nn.ELU(),  
    nn.Conv1d(in_channels=num_channels, out_channels=num_channels,  
              kernel_size=3, stride=2),  
    nn.ELU(),  
    nn.Conv1d(in_channels=num_channels, out_channels=num_channels,  
              kernel_size=3),  
    nn.ELU()  
)
```

## Question 2.2: What other layers are included in the model?

---

- Hint: See Section “The Model: NewsClassifier”, `__init__` and forward functions

# Answer 2.2

---

- `__init__` function

```
self._dropout_p = dropout_p
self.fc1 = nn.Linear(num_channels, hidden_dim)
self.fc2 = nn.Linear(hidden_dim, num_classes)
```

- forward function

```
features = self.convnet(x_embedded)
# average and remove the extra dimension
remaining_size = features.size(dim=2)
features = F.avg_pool1d(features, remaining_size).squeeze(dim=2)
features = F.dropout(features, p=self._dropout_p)
# mlp classifier
intermediate_vector = F.relu(F.dropout(self.fc1(features), p=self._dropout_p))
prediction_vector = self.fc2(intermediate_vector)
if apply_softmax:
    prediction_vector = F.softmax(prediction_vector, dim=1)
```

# Question 2.3: How many channels does each 1D CNN produce as output?

---

- Hints:

- See Section “The Model: NewsClassifier”, `__init__` function
- Section “Initializations”
- Section “Settings and some prep work”

# Answer 2.3

---

- Section “The Model: NewsClassifier”

- `__init__` function

```
self.convnet = nn.Sequential(  
    nn.Conv1d(in_channels=embedding_size,  
              out_channels=num_channels,  
              kernel_size=3),  
    nn.ELU(),  
    nn.Conv1d(in_channels=num_channels,  
              out_channels=num_channels,  
              kernel_size=3, stride=2),  
    nn.ELU(),  
    nn.Conv1d(in_channels=num_channels,  
              out_channels=num_channels,
```

```
              kernel_size=3, stride=2),  
    nn.ELU(),  
    nn.Conv1d(in_channels=num_channels,  
              out_channels=num_channels,  
              kernel_size=3),  
    nn.ELU()  
)
```

- Section “Initializations”

- `classifier = NewsClassifier(..., num_channels=args.num_channels, ...)`

- Section “Settings and some prep work”

- `args = Namespace(..., num_channels=100, ...)`



# Question 2.4: What is difference between the padding introduced in the lecture slides and the padding introduced in the notebook?

---

- Hints:

- Section “The vectorizer”, `vectorize` function
- Section “The dataset”, `__init__` and `__getitem__` functions
- Section “The Vocabulary”, `__init__` function

# Answer 2.4

---

- Section “The vectorizer”, vectorize function

- `indices = [self.title_vocab.begin_seq_index]`
- `indices.extend(self.title_vocab.lookup_token(token) for token in title.split(" "))`
- `indices.append(self.title_vocab.end_seq_index)`
- `if vector_length < 0: vector_length = len(indices)`
- `out_vector = np.zeros(vector_length, dtype=np.int64)`

- `out_vector[:len(indices)] = indices`
- `out_vector[len(indices):] = self.title_vocab.mask_index`

- Section “The dataset”

- `__init__` function
  - `measure_len = lambda context: len(context.split(" "))`
  - `self._max_seq_length = max(map(measure_len, news_df.title)) + 2`
- `__getitem__` function
  - `title_vector = self._vectorizer.vectorize(row.title, self._max_seq_length)`

# Answer 2.4

---

- Section “The Vocabulary”

- `def __init__(self, token_to_idx=None, unk_token="<UNK>", mask_token="<MASK>", begin_seq_token="<BEGIN>", end_seq_token="<END>"):`
  - `self.mask_index = self.add_token(self._mask_token)`
  - `self.unk_index = self.add_token(self._unk_token)`
  - `self.begin_seq_index = self.add_token(self._begin_seq_token)`
  - `self.end_seq_index = self.add_token(self._end_seq_token)`

## Question 2.5: What pooling method is used in the notebook?

---

- Hint: See Section “The Model: NewsClassifier”, forward function

## Answer 2.5

---

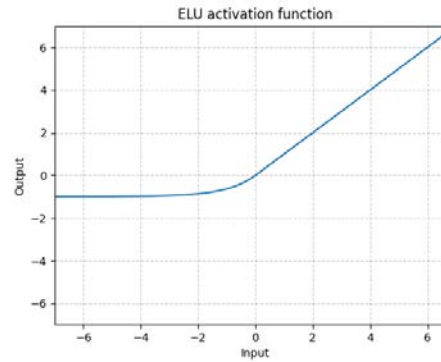
- `features = F.avg_pool1d(features, remaining_size).squeeze(dim=2)`

## Question 2.6: Which regularisation method is used in the notebook?

---

- Hint: See Section “The Model: NewsClassifier”, `__init__` and forward functions

# Answer 2.6



- Section “The Model: NewsClassifier”

- `__init__` function

```
self.convnet = nn.Sequential(
    nn.Conv1d(in_channels=embedding_size,
              out_channels=num_channels,
              kernel_size=3),
    nn.ELU(),
    nn.Conv1d(in_channels=num_channels,
              out_channels=num_channels,
              kernel_size=3, stride=2),
    nn.ELU(),
    nn.Conv1d(in_channels=num_channels,
              out_channels=num_channels,
```

```
              kernel_size=3, stride=2),
    nn.ELU(),
    nn.Conv1d(in_channels=num_channels,
              out_channels=num_channels,
              kernel_size=3),
    nn.ELU()
)
```

- forward function

- `features = F.dropout(features, p=self._dropout_p)`
- `intermediate_vector = F.relu(F.dropout(self.fc1(features), p=self._dropout_p))`

# Answer 2: Summary

---

- 4 layers of 1D CNN
- 4 non-linear layers; 2 fully connected layers; 1 avg pooling layer
- 100 output channels for each 1D CNN
- padding with special token (mask\_index) is used in notebook, while padding of 0.0 is used in lecture slides
- average pooling method
- dropout regularization