

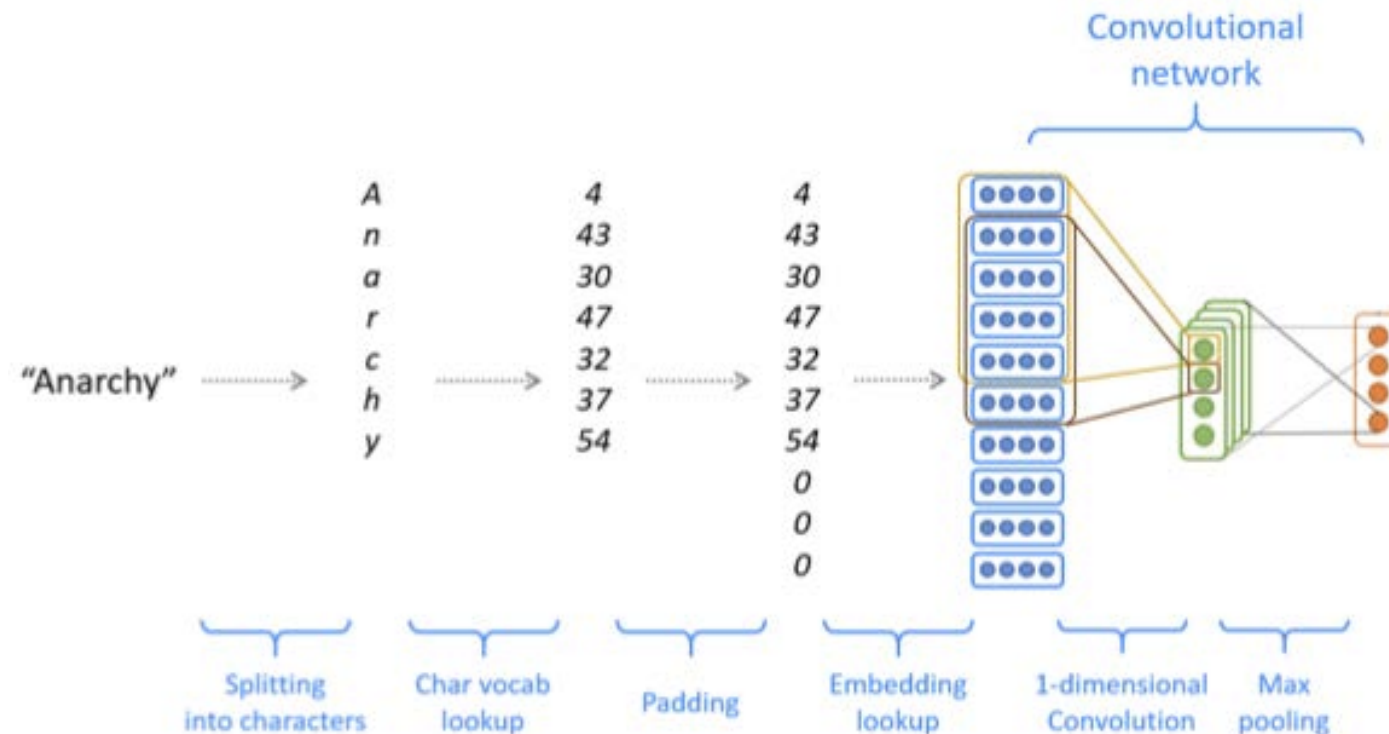
Deep Neural Networks for Natural Language Processing (AI6127)

JUNG-JAE KIM

TUTORIAL 7: SUBWORD MODELS

Question 1: Modify the CNN model of document classification of Tutorial 4 to use character embeddings of “Character-aware neural language models”

- Figure 1. Character-based convolutional encoder, which ultimately produces a word embedding of length e_{word} .



Question 1: Modify the CNN model to use character embeddings of the method “Character-aware neural language models”

- **Convert word to character indices:**
- Assume we have a pre-defined ‘vocabulary’ of characters (e.g. all lowercase letters, uppercase letters, numbers and punctuations).
- By looking up the index of each character, we can represent the length- l word x as a vector of integers: $\mathbf{x} = [c_1, c_2, \dots, c_l] \in \mathbb{Z}^l$
 - where each is c_i an integer index into the character vocabulary.

Question 1: Padding and embedding lookup

- Using a special <PAD> ‘character’, we pad (or truncate) every word so that it has length m_{word} (pre-defined hyper parameter, representing maximum word length):
 - $\mathbf{x}_{padded} = [c_1, c_2, \dots, c_{m_{word}}] \in \mathbb{Z}^{m_{word}}$
- For each of these characters c_i , we lookup a dense character embedding (which has shape e_{char} ; $e_{char} = 50$). This yields a tensor \mathbf{x}_{emb} :
 - $\mathbf{x}_{emb} = \text{CharEmbedding}(\mathbf{x}_{padded}) \in \mathbb{R}^{m_{word} \times e_{char}}$
- We will reshape \mathbf{x}_{emb} to obtain $\mathbf{x}_{reshaped} \in \mathbb{R}^{e_{char} \times m_{word}}$ before feeding into the convolutional network of the Convolutional network below.
 - Necessary because PyTorch Conv1D performs only on the last dimension of the input

m_{word} : max. word len.
 e_{char} : char. emb. size
 k : kernel size
 f : out. channel size

Question 1: Convolutional network

- To combine these character embeddings, we'll use 1-dimensional convolutions.
- The convolutional layer has two hyper-parameters:
 - the kernel size k (also called window size; $k=5$), which dictates the size of the window used to compute features, and
 - the number of filters f (also called number of output features or number of output channels)
 - Assume no padding is applied and the stride is 1
- The convolutional layer has a weight matrix $\mathbf{W} \in \mathbb{R}^{f \times e_{char} \times k}$ and a bias vector $\mathbf{b} \in \mathbb{R}^f$.

Question 1: Modify the model to use character embeddings of the method Character-aware neural language models

m_{word} : max. word len.
 e_{char} : char. emb. size
 k : kernel size
 f : out. channel size

- To compute the i^{th} output feature (where $i \in \{1, \dots, f\}$) for the t^{th} window of the input, the convolution operation is performed between the input window $(\mathbf{x}_{reshaped})_{[:,t:t+k-1]} \in \mathbb{R}^{e_{char} \times k}$ and the weights $\mathbf{W}_{[i,:,:]} \in \mathbb{R}^{e_{char} \times k}$, and the bias term $b_i \in \mathbb{R}$ is added:
 - $(\mathbf{x}_{conv})_{i,t} = \text{sum} \left(\mathbf{W}_{[i,:,:]} \odot (\mathbf{x}_{reshaped})_{[:,t:t+k-1]} \right) + b_i \in \mathbb{R}$
- where \odot is element-wise multiplication of two matrices with the same shape and sum is the sum of all the elements in the matrix. This operation is performed for every feature i and every window t , where $t \in \{1, \dots, m_{word} - k + 1\}$. Overall this produces output \mathbf{x}_{conv} :
 - $\mathbf{x}_{conv} = \text{Conv1D}(\mathbf{x}_{reshaped}) \in \mathbb{R}^{f \times (m_{word} - k + 1)}$

Question 1: Modify the model to use character embeddings of the method Character-aware neural language models

- For our application, we'll set f to be equal to e_{word} , the size of the final word embedding for word x (the rightmost vector in Figure 1).

Therefore,

- $\mathbf{x}_{conv} \in \mathbb{R}^{e_{word} \times (m_{word} - k + 1)}$

- Finally, we apply the ReLU function to \mathbf{x}_{conv} , then use max-pooling to reduce this to a single vector $\mathbf{x}_{conv-out} \in \mathbb{R}^{e_{word}}$, which is the final output of the Convolutional Network:

- $\mathbf{x}_{conv-out} = \text{MaxPool}(\text{ReLU}(\mathbf{x}_{conv})) \in \mathbb{R}^{e_{word}}$

m_{word} : max. word len.

e_{char} : char. emb. size

k : kernel size

f : out. channel size

e_{word} : word emb. size

Question 1: Modify the model to use character embeddings of the method Character-aware neural language models

- Here, MaxPool simply takes the maximum across the second dimension.
- Given a matrix $M \in \mathbb{R}^{a \times b}$, then $\text{MaxPool}(M) \in \mathbb{R}^a$ with $\text{MaxPool}(M)_i = \max_{1 \leq j \leq b} M_{ij}$ for $i \in \{1, \dots, a\}$
- This output is fed into the Convolutional Network of Q1.

Hands-on

Answer 1: Convert word to character indices

The Vocabulary

```
class CharacterSequenceVocabulary(Vocabulary):
    def __init__(self, token_to_idx=None, pad_token="<PAD>"):
        super(CharacterSequenceVocabulary, self).__init__(token_to_idx)
        self._pad_token = pad_token
        self.pad_index = self.add_token(self._pad_token)
    def to_serializable(self):
        contents = super(CharacterSequenceVocabulary, self).to_serializable()
        contents.update({'pad_token': self._pad_token})
        return contents
    def lookup_token(self, token):
        return self._token_to_idx[token]
```

Answer 1: Convert word to character indices

The Vectorizer

```
class NewsVectorizer(object):
    def __init__(self, title_char_vocab, category_vocab):
        self.title_char_vocab = title_char_vocab
        self.category_vocab = category_vocab

    def vectorize(self, title, max_seq_length,
max_word_length):
        words = title.split(" ")
        # cut off title words after max length
        if len(words) > max_seq_length:
            words = words[:max_seq_length]
        out_vectors = []
        for word in words:
            chars = list(word)
            # cut off word characters after max length
            if len(chars) > max_word_length:
                chars = chars[:max_word_length]
            # retrieve character embeddings
```

```
        char_indices =
            [self.title_char_vocab.lookup_token(token)
             for char in chars]
        out_vector = np.zeros(max_word_length,
dtype=np.int64)
        out_vector[:len(char_indices)] = char_indices
        # fill up with <PAD> embeddings
        if len(char_indices) < max_word_length:
            out_vector[len(char_indices):] =
                self.title_char_vocab.pad_index
        out_vectors.append(out_vector)
        if len(words) < max_seq_length:
            null_word_emb =
np.array([self.title_char_vocab.pad_index] *
max_word_length, dtype=np.int64)
            for _ in range(max_seq_length - len(words)):
                out_vectors.append(null_word_emb)
        out_vectors = np.array(out_vectors, dtype=np.int64)
        return out_vectors
```

Answer 1: Convert word to character indices

The Vectorizer

```
def from_dataframe(cls, news_df):
    ...
    title_char_vocab =
        CharacterSequenceVocabulary()
    for title in news_df.title:
        for token in title.split(" "):
            title_char_vocab.\
                add_many(list(token))

def from_serializable(cls, contents):
    title_char_vocab =
        CharacterSequenceVocabulary.from_seriali
        zable(contents['title_char_vocab'])
    category_vocab =
        Vocabulary.from_serializable(contents['cat
        egory_vocab'])
    return cls(
        title_char_vocab=title_char_vocab,
        category_vocab=category_vocab)

def to_serializable(self):
    return {'title_char_vocab':
        self.title_char_vocab.to_serializable(),
        'category_vocab':
        self.category_vocab.to_serializable()}
```

Answer 1: Convert word to character indices

The Dataset

```
class NewsDataset(Dataset):
    def __init__(self, news_df, vectorizer):
        ...
        self._max_seq_length = max(map(measure_len, news_df.title))
        self._max_word_length = 0
        for title in news_df.title:
            for token in title.split(" "):
                if len(token) > self._max_word_length:
                    self._max_word_length = len(token)
    def __getitem__(self, index):
        ...
        title_vector = self._vectorizer.vectorize(row.title, self._max_seq_length,
            self._max_word_length)
        ...
```

m_{word} : max. word len.

Answer 1: Convert word to character indices

The Model: NewsClassifier

```
class NewsClassifier(nn.Module):
    def __init__(self, char_embedding_size, word_embedding_size, char_num_embeddings,
                  word_num_channels, char_kernel_size, hidden_dim, num_classes, dropout_p,
                  char_pretrained_embeddings=None, padding_idx=0):
        super(NewsClassifier, self).__init__()
        if char_pretrained_embeddings is None:
            self.char_emb = nn.Embedding(embedding_dim=char_embedding_size, echar: char. emb. size
                                         num_embeddings=char_num_embeddings,
                                         padding_idx=padding_idx)
        else:
            char_pretrained_embeddings =
                torch.from_numpy(char_pretrained_embeddings).float()
            self.char_emb = nn.Embedding(embedding_dim=char_embedding_size,
                                         num_embeddings=char_num_embeddings,
                                         padding_idx=padding_idx, _weight=char_pretrained_embeddings)
```

Answer 1: Convert word to character indices

The Model: NewsClassifier

```
def forward(self, x_in, apply_softmax=False):  
    # x_in: (batch_size, max_seq_size, max_word_size)  
    # x_emb: (batch_size, max_seq_size, max_word_size, char_embedding_size)  
    x_emb = self.char_emb(x_in)  
    batch_size = x_emb.size(dim=0)  
    max_seq_size = x_emb.size(dim=1)  
    max_word_size = x_emb.size(dim=2)  
    char_embedding_size = x_emb.size(dim=3)  
    # x_resaped: (batch_size * max_seq_size, char_embedding_size, max_word_size)  
    x_resaped = x_emb.view(batch_size * max_seq_size, max_word_size,  
                           char_embedding_size).permute(0, 2, 1)
```

Answer 1: Convolutional network

The Model: NewsClassifier

```
def __init__(...):
```

```
...
```

```
self.char_convnet = nn.Sequential(  
    nn.Conv1d(in_channels=char_embedding_size,  
              out_channels=word_embedding_size,  
              kernel_size=char_kernel_size),  
    nn.ReLU())
```

m_{word} : max. word len.
 e_{char} : char. emb. size
k: kernel size
f: out. channel size

```
def forward(self, x_in, apply_softmax=False):
```

```
# x_conv: (batch_size * max_seq_size, word_embedding_size, max_word_size-char_kernel_size+1)
```

```
x_conv = self.char_convnet(x_resaped)
```

```
# x_conv_out: (batch_size * max_seq_size, word_embedding_size)
```

```
word_embedding_size = x_conv.size(dim=1)
```

```
remaining_size = x_conv.size(dim=2)
```

```
x_conv_out = F.max_pool1d(x_conv, remaining_size).squeeze(dim=2)
```


Answer 1: This output is fed into the Convolutional Network of Q1

The Model: NewsClassifier

```
def forward(self, x_in, apply_softmax=False):
```

```
...
```

```
    features = self.word_convnet(x_conv_out.view(
        batch_size, max_seq_size, word_embedding_size).permute(0, 2, 1))
```

Answer 1: Training

Settings and some prep work

```
args = Namespace(...,  
    word_embedding_size=100, char_embedding_size=50,  
    char_kernel_size=5, word_num_channels=100, ...)
```

Initializations

```
args.use_glove = False  
classifier = NewsClassifier(char_embedding_size=args.char_embedding_size,  
    word_embedding_size=args.word_embedding_size,  
    char_num_embeddings=len(vectorizer.title_char_vocab),  
    word_num_channels=args.word_num_channels,  
    char_kernel_size=args.char_kernel_size,  
    char_pretrained_embeddings=embeddings, ...)
```